

Manuel Logistic Regression

May 21, 2024

```
[56]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
[63]: def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def sgd(X, y, batch_size=10, learning_rate=0.1, max_iterations=100):
    np.random.seed(0)
    num_features = X.shape[1]
    w = np.random.randn(num_features)
    n = len(y)
    for iteration in range(max_iterations):
        indices = np.random.permutation(n)
        for start in range(0, n, batch_size):
            end = min(start + batch_size, n)
            batch_indices = indices[start:end]
            X_batch = X[batch_indices]
            y_batch = y[batch_indices]

            z = np.dot(X_batch, w)
            y_hat = sigmoid(z)

            gradient = np.dot(X_batch.T, (y_hat - y_batch))

            w -= learning_rate * gradient

    return w
```

```
[77]: from sklearn.datasets import load_breast_cancer
data = load_breast_cancer()
X = data.data
y = data.target
```

```
[108]: df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target
# Check for missing values
```

```

numMissing = df.isnull().sum().sum()
print(numMissing)
# Standardize values
print(df.info())
scaler = StandardScaler()
X = scaler.fit_transform(X)
# Split
X_train_val, X_test, y_train_val, y_test = train_test_split(X, y, test_size=0.
↪2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val, ↪
↪test_size=0.25, random_state=42)

```

0

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	mean radius	569 non-null	float64
1	mean texture	569 non-null	float64
2	mean perimeter	569 non-null	float64
3	mean area	569 non-null	float64
4	mean smoothness	569 non-null	float64
5	mean compactness	569 non-null	float64
6	mean concavity	569 non-null	float64
7	mean concave points	569 non-null	float64
8	mean symmetry	569 non-null	float64
9	mean fractal dimension	569 non-null	float64
10	radius error	569 non-null	float64
11	texture error	569 non-null	float64
12	perimeter error	569 non-null	float64
13	area error	569 non-null	float64
14	smoothness error	569 non-null	float64
15	compactness error	569 non-null	float64
16	concavity error	569 non-null	float64
17	concave points error	569 non-null	float64
18	symmetry error	569 non-null	float64
19	fractal dimension error	569 non-null	float64
20	worst radius	569 non-null	float64
21	worst texture	569 non-null	float64
22	worst perimeter	569 non-null	float64
23	worst area	569 non-null	float64
24	worst smoothness	569 non-null	float64
25	worst compactness	569 non-null	float64
26	worst concavity	569 non-null	float64
27	worst concave points	569 non-null	float64
28	worst symmetry	569 non-null	float64
29	worst fractal dimension	569 non-null	float64

```
30 target          569 non-null    int64
dtypes: float64(30), int64(1)
memory usage: 137.9 KB
None
```

```
[81]: from collections import Counter
      class_counts = Counter(y_train)
      print("Size of classes in training set: ", class_counts)
```

```
Size of classes in training set: Counter({1: 216, 0: 125})
```

```
[105]: weights = sgd(X_train, y_train, batch_size=10, learning_rate=0.01,
      ↪max_iterations=100)
      def predict(X, weights):
          z = np.dot(X, weights)
          return sigmoid(z)
```

```
[106]: y_pred = (predict(X_test, weights) >= 0.5).astype(int)
      accuracy = accuracy_score(y_test, y_pred)
      precision = precision_score(y_test, y_pred)
      recall = recall_score(y_test, y_pred)
      f1 = f1_score(y_test, y_pred)

      print(f"Accuracy: {accuracy:.4f}, Precision: {precision:.4f}, Recall: {recall:.4f}, F1-score: {f1:.4f}")
```

```
Accuracy: 0.9737, Precision: 0.9857, Recall: 0.9718, F1-score: 0.9787
```

Summary:

Based on these findings, it is obvious that the model does a great job with over a 0.97 in accuracy, precision, recall, and f1-score. I separated my code based on each part and my code is pretty straightforward. The linear regression uses all the predictors which are all floats to predict either a one or a zero.